


CG1112

Engineering Principles and Practices II for CEG

Week 7 Studio 2

Communication Protocols




Introduction

- **In the previous studio you learnt how to establish communications with your Arduino using serial communications.**
- **In this lecture we will look at “communications protocols”**
 - A “protocol” is an agreement between machines on how to represent data.
- **In particular we will be looking at application layer protocols.**
 - In the previous studio we saw:
 - ✓ **Physical layer protocols – Agreement on voltage levels, number of wires to use, how wires are to be connected, etc.**
 - ✓ **Link layer protocols – How bits are placed together into units called “frames”, what the layout of each frame should be, etc.**
 - **Application layer protocols:**
 - ✓ **How we arrange instructions and data within a data structure for controlling behaviors of remote systems.**

Three things to do:

- **Decide on physical connection:**
 - Easy enough; we will connect the Arduino to the Pi using USB.
- **Decide on bit-level protocol:**
 - Decide baud rate.
 - Decide data length.
 - Decide # of parity bits.
 - Decide # of stop bits.
 - We will use 9600 8N1 as standard..



Serial Communication and Protocol Design

BUILDING A PROTOCOL

Assign an ID to each device

- You need to be able to identify sensors (actuators) to read from (send data to).

Device ID	Device
0	Sonar 1
1	Sonar 2
2	Touch Sensor 1
3	Touch Sensor 2
4	Buzzer
5	Tactile feedback motor
...	...

Create Packet Types


- So both sides know what sort of packets are being sent (and the appropriate response)

Packet Type	Packet Code
ACK	0
NAK	1
Hello	2
Read	3
Write	4
Data Response	5
...	...

Getting Data from the Arduino


- **Your Arduino may have data to send back to the Raspberry Pi (RPi):**
 - Compass readings.
 - Number of turns the wheels have made.
 - Etc?
- **You have two choices:**
 - Arduino can periodically send back data in “heartbeat packets”.
 - ✓ These are packets sent back say once per second, containing data and status information.
 - RPi can poll for data.
 - ✓ RPi sends a request for data.
 - ✓ Arduino replies.

Periodic Push By Arduino



- **Arduino sends data whenever it is available.**
- **Often implemented as a “heartbeat” packet.**
- **RPi monitors and buffers data as it comes in.**
 - +Arduino sends data whenever it is available.
 - RPi needs to buffer incoming data.



Periodic Poll by RPi



- **Arduino waits for poll packets from RPi.**
- **RPi requests data when it needs it.**
 - +RPi decides when it needs the data and sends poll packet.
 - If RPi doesn't poll often enough, may lose data on Arduino (Arduino has only 2K of RAM)

Commanding the Arduino

- **You can use a variation of the “poll” packet to send commands to the Arduino:**
 - E.g. tell Arduino to move Vincent forward by 1 meter at 50% speed.
- **The Arduino responds immediately with an ACK packet.**
- **The RPi can either:**
 - Poll the Arduino continuously to check when it has finished the command, OR
 - The Arduino can send back a packet to inform the RPi that it is done.



Finding Checksums

- **Checksums are used to check that data is received correctly.**
- **Sender side:**
 - Compute checksum
$$\text{checksum} = b1 \text{ XOR } b2 \text{ XOR } b3 \text{ XOR } b4 \text{ XOR } \dots$$
 - Attach to end of packet.
- **Receiver side:**
 - Compute checksum using data in packet, except checksum.
 - Compare against attached checksum.
 - If equal, reply with ACK, else reply with NAK.

Finding Checksums

- **For the most part:**
 - Serial comms is reliable.
 - Hence we don't normally compute checksums (or send ACK for that matter).
- **However:**
 - Your set up is not going to be perfect. (headers and pins! C'mon)
 - If you are sending relatively large amounts of data, higher chance of errors.

Serializing Structures

- **The easiest way to implement the protocol packets is as structures.**
- **Example for our command packet:**

```
typedef tc
{
    int command;
    int speed;
    int distanceInMeters;
} TCommand;
```

- **Serializing: Converting a structure into a stream of bytes, because serial devices can only deal with streams of bytes.**
 - Get a pointer to the structure.
 - Copy into an array of char.
 - May want to include information on packet length and checksum.

Serializing Structures

```
typedef struct con
{
    unsigned char devCode;
    double maxValue;
    double minValue;

} TConfigPacket;

void sendConfig()
{
    TConfigPacket cfg;
    char buffer[64];
    cfg.devCode=deviceCode;
    cfg.minValue = minValue;
    cfg.maxValue = maxValue;
    unsigned len = serialize(buffer, &cfg, sizeof(cfg));
}
```

Serializing Structures

```
unsigned int serialize(char *buf, void *p, size_t size)
{
    char checksum = 0;
    buf[0]=size;
    memcpy(buf+1, p, size);
    for(int i=1; i<=size; i++)
    {
        checksum ^= buf[i];
    }
    buf[size+1]=checksum;
    return size+2;
}

void sendSerialData(char *buffer, int len)
{
    for(int i=0; i<len; i++)
        Serial1.write(buffer[i]);
}
```

Deserializing Structures

- **Deserialize: Convert a stream of bytes back to structures.**
 - Get a pointer to the structure.
 - Copy buffer of bytes to that pointer:
 - ✓ **May need to remove packet length and compute checksums first.**

Deserializing Structures

```
Void readConfig()  
{  
    char buffer[MAX_BUF_LEN];  
    int len;  
    TConfigPacket cfg;  
  
    readSerial(buffer, &len);  
    deserialize(&cfg, buffer);  
  
    // Process cfg.  
    ...  
}
```

Deserializing Structures

```
unsigned int deserialize(void *p, char *buf)
{
    size_t size = buf[0];
    char checksum = 0;

    for(int i=1; i<=size; i++)
        checksum ^= buf[i];

    if(checksum == buf[size+1])
    {
        memcpy(p, buf+1, size);
        return PACKET_OK;
    }
    else
    {
        printf("CHECKSUM ERROR\n");
        return PACKET_BAD_CHECKSUM;
    }
}
```

Serializing / Deserializing Structures


- **Our serializing / deserializing algorithms may not work in the real world because of two complications:**
 - Endianness
 - Differing Data Types
- **In addition:**
 - The representation shown earlier is not efficient. We would use a separate data structure to store the data to be serialized, the checksum and the size of the data.
 - See `serialize.zip` for an example.

Complication #1


Endianness

How to store multibyte word (object > 1 byte), 2 general schemes, Eg: **0x87654321** (4 byte int)

- **Big Endian**: higher order bytes at lower addresses



- **Little Endian**: lower order bytes at lower addresses



x86: little endian ; Sparc: big endian ; powerpc: configurable

Complication #1

Endianness

- **As it turns out:**
 - The Arduino is little endian.
 - The Pi is little endian.
- **Our job here is done. ;)**
- **Note though:**
 - If your destination machine has a different endianness, our algorithm will not work.
 - Serializing / Deserializing will be considerably more troublesome.
 - ✓ **Convert all data to a standard endianness.**
 - ✓ **At the destination, convert it back to the native endianness.**
 - In such cases consider using established libraries like MAVLINK and protobuf.



Complication #2

Different Data Sizes

- **We want to send over a data structure consisting of:**
 - A single character, initialized to 'a'.
 - An integer, initialized to 1.
- **As it turns out, the following code segments do not work. The sample printout on the Pi is shown on the next page:**

Arduino Side

```
void setup() {  
    // put your setup code here, to run once:  
  
    Serial.begin(115200);  
}  
  
typedef struct tm  
{  
    int x;  
    char c;  
} TTest;  
  
void loop() {  
    TTest data;  
  
    data.x=1;  
    data.c='a';  
  
    // put your main code here, to run repeatedly:  
  
    Serial.write((char *) &data, sizeof(TTest));  
    delay(500);  
}
```

RPi Side

```
#include <stdio.h>  
#include <termios.h>  
#include <string.h>  
#include <stdint.h>  
#include "serial.h"  
  
typedef struct tm  
{  
    int x;  
    char c;  
} TTest;  
  
int main()  
{  
    char buffer[MAX_BUFFER_LEN];  
    int n;  
  
    memset(buffer, 0, MAX_BUFFER_LEN);  
  
    startSerial("/dev/ttyACM0", B115200, 8, 'N', 1, 5);  
    while(1)  
    {  
        n = serialRead(buffer);  
        TTest *data = (TTest *) buffer;  
        printf("Read %d bytes: x is %d, c is %c\n", n, data->x, data->c);  
    }  
}
```

Complication #2

Different Data Sizes

Intended Output (RPi)

```
ATTEMPTING TO CONNECT TO SERIAL. ATTEMPT # 1 of 5.  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a  
Read 3 bytes: x is 1, c is a
```

Actual Output (RPi)


```
ATTEMPTING TO CONNECT TO SERIAL. ATTEMPT # 1 of 5.  
Read 6 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is  
Read 3 bytes: x is 23134209, c is
```

- We can see that the output is obviously wrong. BUT WHY?

Complication #2

Different Data Sizes


- **Turns out:**
 - Ints on the Arduino are 16 bits wide, but ints on the RPi are 32 bits wide!
 - The diagram below shows the consequences:



What the
Arduino sent

Our integer
 $x=1$

Our character
 $c='a'$



What the Pi
THINKS the
Arduino sent

Where the Pi thinks x is.
We end up with garbage

Where the Pi
thinks c is.

Complication #2

Different Data Sizes

- **Fortunately there are two solutions:**
 - Switch to Arduino Due, which has a 32-bit ARM architecture like the Pi.
 - ✓ However the hardware registers and how we access them is completely different.
 - ✓ Despair, fail CG1112, drop out and spend the rest of our lives scraping dead animals off the streets.
 - Use standardized integer types
 - Replace:
 - int with int32_t*
 - unsigned int with uint32_t*
 - long with int64_t*
 - unsigned long with uint64_t*
 - ✓ **Must remember to #include <stdint.h>**
 - Floats are not affected.
 - ✓ **Both the Arduino and Pi use 32-bit IEEE 754 format in little endian.**

Complication #2

Different Data Sizes

- Our corrected code (with additional integer y and a floating point z thrown in!):

Arduino Side

```
void setup() {
    // put your setup code here, to run once:

    Serial.begin(115200);
}

typedef struct tm
{
    int32_t x;
    int32_t y;
    float z;
    char c;
} TTest;

void loop() {
    TTest data;

    data.x=1;
    data.y=2;
    data.z = 3.141592654;
    data.c='a';

    // put your main code here, to run repeatedly:

    Serial.write((char *) &data, sizeof(TTest));
    delay(500);
}
```

RPi Side

```
#include <stdio.h>
#include <termios.h>
#include <string.h>
#include <stdint.h>
#include "serial.h"

typedef struct tm
{
    int32_t x;
    int32_t y;
    float z;
    char c;
} TTest;

int main()
{
    char buffer[MAX_BUFFER_LEN];
    int n;

    memset(buffer, 0, MAX_BUFFER_LEN);

    startSerial("/dev/ttyACM0", B115200, 8, 'N', 1, 5);
    while(1)
    {
        n = serialRead(buffer);
        TTest *data = (TTest *) buffer;
        printf("Read %d bytes: x is %d, y is %d, z is %f, c is %c\n",
            n, data->x, data->y, data->z, data->c);
    }
}
```

Complication #2

Different Data Sizes

Output (RPi side)

[illegible]

Bonus Complication

- **One bonus complication:**

- The serial receive doesn't guarantee that you will receive the entire packet at one time.
- You may receive fragments, which can cause errors as shown below:

```
Read 13 bytes: x is 1, y is 2, z is 3.141593, c is a
Read 13 bytes: x is 1, y is 2, z is 3.141593, c is a
Read 1 bytes: x is 1, y is 2, z is 3.141593, c is a
Read 12 bytes: x is 33554432, y is -620756992, z is 221689955540103331840.000000, c is a
Read 3 bytes: x is 33554433, y is -620756992, z is 221689955540103331840.000000, c is a
Read 10 bytes: x is 512, y is 1225775872, z is 221798903948276400128.000000, c is a
```

- **Solution:**

✓ **Test number of bytes received. If it is == size of data structure, accept. Otherwise buffer and add in subsequent bytes that arrive.**

Bonus Complication

- **We do two things:**
 - We pad the TTest structure so that its size in bytes is divisible by 4
 - This is to match what the compiler does on the Pi side:
 - ✓ We currently have 2 int32_t, 1 float and 1 char.
 - ✓ This totals $4 + 4 + 4 + 1 = 13$ bytes.
 - ✓ We add in a dummy of 3 bytes:

```
typedef struct tm
{
    int32_t x;
    int32_t y;
    float z;
    char c;
    char dummy[3];
} ttest;
```

Bonus Complication

- Now we write the code to re-assemble the packet fragments:

```
typedef enum
{
    PACKET_OK = 0,
    PACKET_INCOMPLETE = 1,
    PACKET_ERROR = 2
} TResult;

TReturn assemble(char *packetBuffer, const char *dataBuffer, int len, size_t dataLen)
{
    static int count=0;

    if(count + len > dataLen)
    {
        count=0;
        return PACKET_ERROR;
    }
    else
    {
        int i;
        for(i=0; i<len; i++)
            packetBuffer[count++]=dataBuffer[i];
    }

    if(count == dataLen)
    {
        count=0;
        return PACKET_OK;
    }
    else
        return PACKET_INCOMPLETE;
}
```

Bonus Complication

- Then we change our main to call assemble:

```
int main()
{
    char buffer[MAX_BUFFER_LEN];
    char dataBuffer[MAX_BUFFER_LEN];
    int n;

    memset(buffer, 0, MAX_BUFFER_LEN);

    startSerial("/dev/ttyACM0", B115200, 8, 'N', 1, 5);
    while(1)
    {
        n = serialRead(buffer);

        TResult result = assemble(dataBuffer, buffer, n, sizeof(TTest));

        if(result == PACKET_OK)
        {
            TTest *data = (TTest *) dataBuffer;
            printf("Read %d bytes: x is %d, y is %d, z is %f, c is %c\n",
                n, data->x, data->y, data->z, data->c);
        }
        else
        {
            if(result == PACKET_ERROR)
                printf("\n\t*** PACKET ERROR ***\n\n");
            else
                printf("Incomplete packet. Expecting %d more bytes.\n", sizeof(TTest) - n);
        }
    }
}
```

Full Serialization / Deserialization Code

- **Please see the `serialize.zip` file for the full serialization/deserialization code.**
- **As this is a rather complicated topic:**
 - Please use the code in this file for your studios.
 - However you **MUST** understand how it works.
 - ✓ In particular we use the **Tcomms** data structure to store the data to be sent, the length of the data, and a checksum.
 - ✓ We also add a magic number. When we receive a packet, we check the magic number to ensure that this is valid packet.